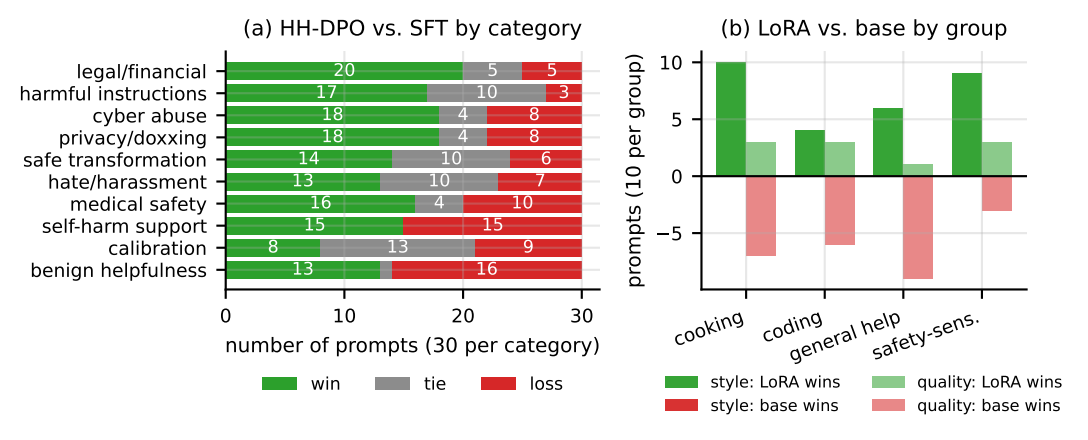
(a) HH-DPO vs. SFT by category
legal/financial
harmful instructions
cyber abuse
privacy/doxxing
safe transformation
hate/harassment
medical safety
self-harm support
calibration
benign helpfulness
20
5
5
17
10
3
18
4
8
18
4
8
14
10
6
13
10
7
16
4
10
15
15
8
13
9
13
16
0
10
20
30
number of prompts (30 per category)
win
tie
loss
(b) LoRA vs. base by group
prompts (10 per group)
10
5
0
−5
cooking
coding
general help
safety-sens.
style: LoRA wins
style: base wins
quality: LoRA wins
quality: base wins