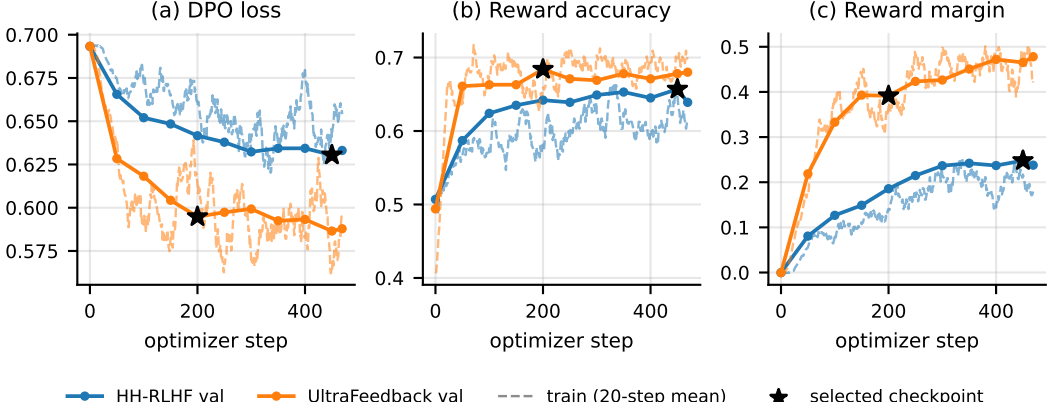
(a) DPO loss
(b) Reward accuracy
(c) Reward margin
0.700
0.675
0.650
0.625
0.600
0.575
0.7
0.6
0.5
0.4
0.5
0.4
0.3
0.2
0.1
0.0
0
200
400
optimizer step
HH-RLHF val
UltraFeedback val
train (20-step mean)
selected checkpoint